



Emily Stolfo

Ruby Engineer at  **mongoDB**, Adjunct Faculty at Columbia
@EmStolfo

**There is no such thing as
thread-safe Ruby code.**

Emily Stolfo



Ruby Engineer at  mongoDB
Adjunct Faculty at Columbia
Expert



THREAD SAFETY FIRST

- 1. Concurrency in Ruby**
- 2. Writing thread-safe code**
- 3. Testing concurrency**

Demo code

```
require 'set'

members = Set.new
threads = []

200.times do |n|
  threads << Thread.new do
    if n % 2 == 0
      members << n
    else
      members.first.nil?
    end
  end
end

threads.each(&:join)
```

Inconsistent bug?

- ✓ **2.0** with **200** threads
- ✓ **JRuby** with **10** threads
- x** **JRuby** with **200** threads

**There are different
Ruby implementations.**

**Different Ruby
implementations have
their own semantics.**

**Threads are like
music.**



GIL



green thread



native (OS) thread

Threads and Ruby

ruby 1.8



Instruments: OS threads

Notes: green threads

Conductor: GIL

Threads and Ruby

ruby 1.9+



Instruments: OS threads

Notes: green threads

Conductor: GIL

Threads and Ruby

JRuby



Instruments: OS threads

Notes: green threads

Different Rubies?
Different semantics.

**Our code
sometimes works by
sheer luck.**

**You're lucky your code
hasn't run on JRuby.**

**You're lucky
there is a GIL.**

Example:

**MongoDB Ruby driver
and Replica Sets**

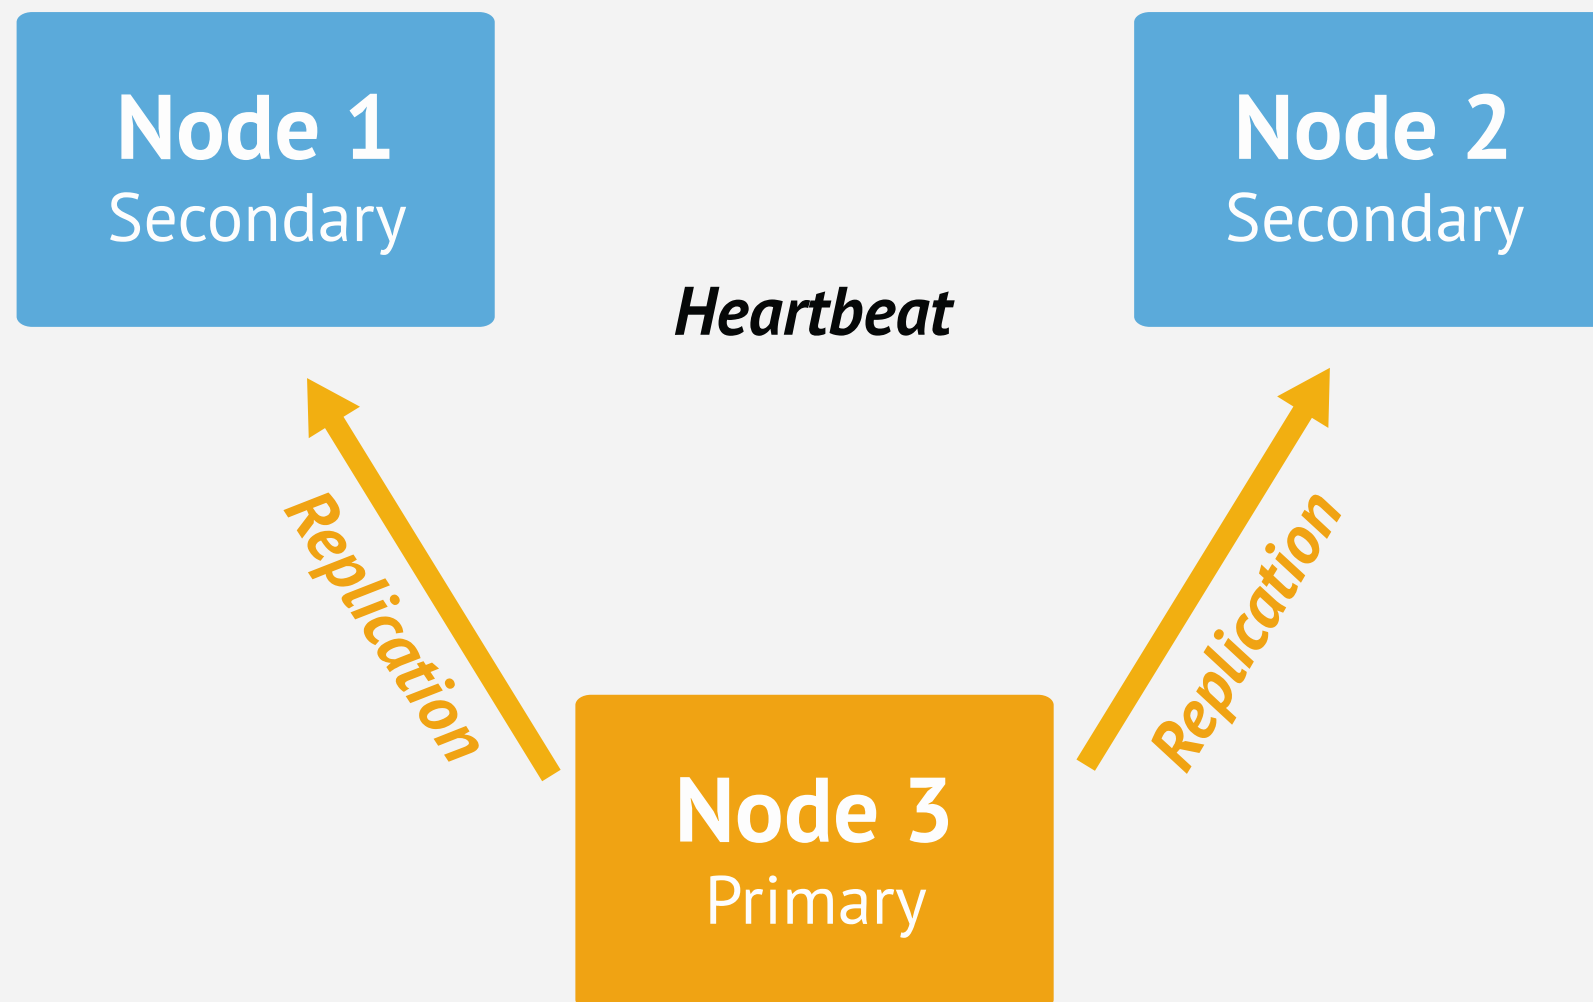
MongoDB Replica Set Creation

Node 1

Node 2

Node 3

MongoDB Replica Set



MongoDB Replica Set Failure

Node 1
Secondary

Node 2
Secondary

No 3

MongoDB Replica Set Recovery



The diagram illustrates a MongoDB Replica Set Recovery scenario. It features a light gray rectangular background. In the upper left, a blue rounded rectangle contains the text 'Node 1' and 'Secondary'. In the upper right, an orange rounded rectangle contains the text 'Node 2' and 'Primary'. In the lower center, there is a gray rounded rectangle. No lines or arrows connect these nodes.

Node 1
Secondary

Node 2
Primary

MongoDB Replica Set Recovered

Node 1
Secondary

Node 2
Primary



Mutable shared state.

Ruby driver

```
class ReplicaSetClient
  def initialize
    @nodes = Set.new
  end

  ...

  def connect_to_nodes
    seed = valid_node
    seed.node_list.each do |host|
      node = Node.new(host)
      @nodes << node if node.connect
    end
  end

  def choose_node(type)
    @nodes.detect { |n| n.type == type }
  end
end
```

Potential concurrency bug

```
class ReplicaSetClient
  def initialize
    @nodes = Set.new
  end

  ...


  def connect_to_nodes
    seed = valid_node
    seed.node_list.each do |host|
      node = Node.new(host)
      @nodes << node if node.connect
    end
  end

  def choose_node(type)
    @nodes.detect { |n| n.type == type }
  end
end
```

(RuntimeError)

"can't add a new key into
hash during iteration"

**We often use
hashes as caches.**



**Hashes and their
derivatives are not
thread-safe in JRuby.**

**How do we write this
code thread-safely?**

WRITING THREAD-SAFE CODE

Shared data:
Avoid across threads.

**If you can't avoid shared
data, at least avoid
shared mutable data.**

**If you can't avoid
shared mutable data,
use concurrency primitives.**

The top 2 concurrency primitives

1. Mutex

```
class ReplicaSetClient
  def initialize
    @nodes = Set.new
    @connect_mutex = Mutex.new
  end

  ...

  def connect_to_nodes
    seed = valid_node
    seed.node_list.each do |host|
      node = Node.new(host)
      @connect_mutex.synchronize do
        @nodes << node if node.connect
      end
    end
  end

  def choose_node(type)
    @connect_mutex.synchronize do
      @nodes.detect { |n| n.type == type }
    end
  end
end
```

1. Mutex

Use to isolate code
that should be
executed by at most
one thread at a time.

```
class ReplicaSetClient
  def initialize
    @nodes = Set.new
    @connect_mutex = Mutex.new
  end

  ...

  def connect_to_nodes
    seed = valid_node
    get_node_from_seed(seed) rescue nil
    node = Node.new(host)

    @connect_mutex.synchronize do
      @nodes << node if node.connect
    end
  end

  def choose_node(type)
    @connect_mutex.synchronize do
      @nodes.detect { |n| n.type == type }
    end
  end
end
```

shared
replica
set
state
update



A mutex is not magic.

**Avoid locking
around I/O.**

1. Mutex

```
class ReplicaSetClient
  def initialize
    @nodes = Set.new
    @connect_mutex = Mutex.new
  end

  ...

  def connect_to_nodes
    seed = valid_node
    seed.node_list.each do |host|
      node = Node.new(host)
      @connect_mutex.synchronize do
        @nodes << node if node.connect
      end
    end
  end

  def choose_node(type)
    @connect_mutex.synchronize do
      @nodes.detect { |n| n.type == type }
    end
  end
end
```

← network
I/O

1. Mutex

```
class ReplicaSetClient
  def initialize
    @nodes = Set.new
    @connect_mutex = Mutex.new
  end
```

```
...
```

```
  def connect_to_nodes
    seed = valid_node
    seed.node_list.each do |host|
      node = Node.new(host)
```

```
      @connect_mutex.synchronize do
```

```
        @nodes << node
```

```
      end if node.connect
```

```
    end
```

```
  end
```

```
  def choose_node(type)
```

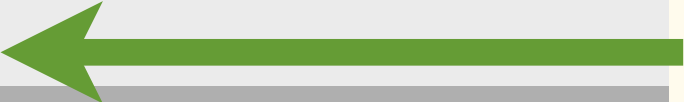
```
    @connect_mutex.synchronize do
```

```
      @nodes.detect { |n| n.type == type }
```

```
    end
```

```
  end
```


```
end
```



network
I/O


Consider resources.

Ex: Thundering herd



Thundering herd

n threads are woken up after waiting on something, but only 1 can continue. Waste of system resources to wake up n threads.



2. Condition Variable

Use to communicate between threads.


Condition Variable

```
class Pool
  def initialize
    @socket_available = ConditionVariable.new
    ...
  end

  def checkout
    loop do
      @lock.synchronize do
        if @available_sockets.size > 0
          socket = @available_sockets.next
          return socket
        end
      end
      @socket_available.wait(@lock)
    end
  end

  def checkin(socket)
    @lock.synchronize do
      @available_sockets << socket
      @socket_available.broadcast
    end
  end
end
```

Why is this
a waste of
system
resources?



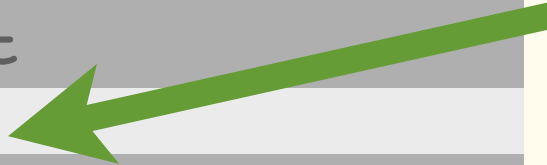
Condition Variable

```
class Pool
  def initialize
    @socket_available = ConditionVariable.new
    ...
  end

  def checkout
    loop do
      @lock.synchronize do
        if @available_sockets.size > 0
          socket = @available_sockets.next
          return socket
        end
      end
      @socket_available.wait(@lock)
    end
  end

  def checkin(socket)
    @lock.synchronize do
      @available_sockets << socket
      @socket_available.signal
    end
  end
end
```

Only one
thread
can
continue



TESTING CONCURRENCY

Testing

- 1. Test with different implementations.**
- 2. Test with a ton of threads.**
- 3. Use patterns for precision.**

**Use synchronization
methods if you need
more precision.**

**ex: rendezvous / barrier
pattern**


Concurrency in Ruby



Know your implementations.



Know your concurrency primitives.



Know your code.

“thread-safe JRuby 1.7.4 code”

Thanks



Emily Stolfo

@EmStolfo